

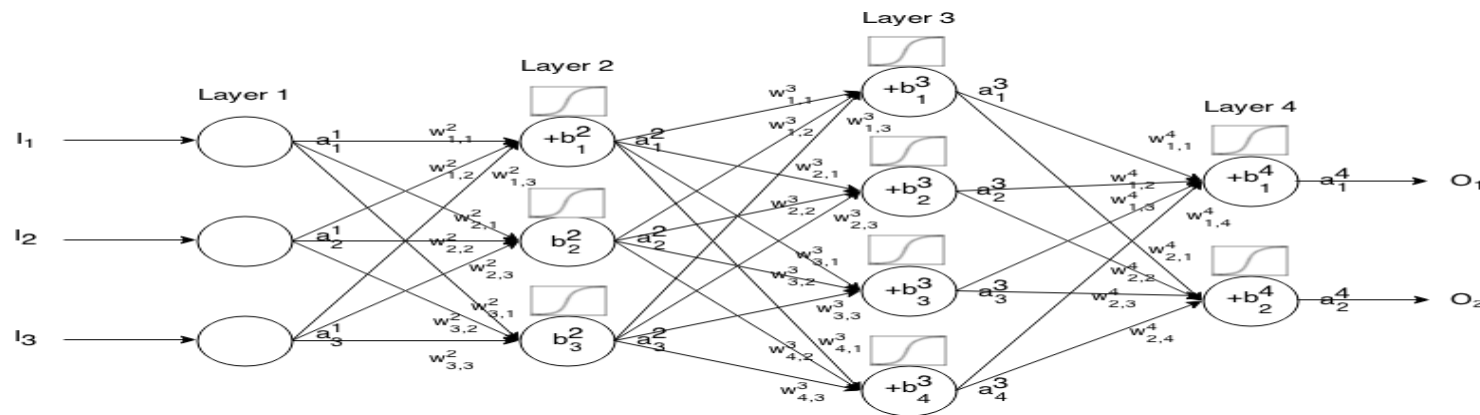
# Artificial Intelligence and Machine Learning. 6CS012.

Lecture-5:Multi layer Neural Network.

**Siman Giri.**


# A. Motivation for Multi Layer-Neural Network.

What are Multi-layer Neural Network?




# Story So Far....

- MCP was first simple computational Model that emulates Human Neuron behavior.




- Perceptron were computationally better representations of Human Neuron.


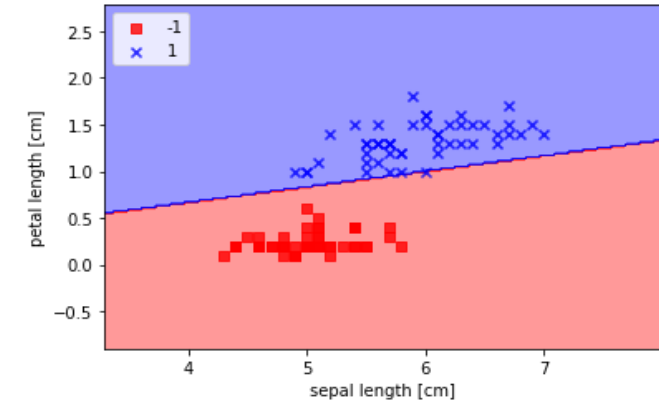


Perceptron rule.

l.




- Linear Separability-Decision Boundry.




- Multi-Layer of Perceptron and Theory of Universal Approximations.

# What are Multi-layer Neural Networks?


Single Perceptron



One-Hidden Layer




Two-Hidden Layer



# 1. Activation Functions.

# 1.1 What are Activation Functions?



- Activation Functions introduce **non-linearity** to the **output** of neurons.
- The activation function does the **non-linear transformation to the input**, making it **capable to learn and perform more complex tasks**.
- The activation function should be **differentiable** or the concept of updating weights (**Backpropagation**) fails, which is the core idea of deep learning.


## 1.2 Some Common Non Linear Activation Functions.

### Sigmoid

- Mathematically;

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

- Squashes real numbers to range between [0,1].



### Pros and Cons.

- the sigmoid non-linearity has recently fallen out of favor and it is rarely ever used.
  - Logistic sigmoid can cause a neural network to get “stuck” during training. This is because, if a strongly-negative input is provided to the logistic sigmoid, it outputs a value, which is very near to zero.
  - Because of this behavior, updating weights will be slow and they are less regularly updated. In simple terms, weights which are updated through back-propagation will be quite slow.


## 1.2 Some Common Non Linear Activation Functions.

### Tanh(Hyperbolic Tangent Function)

- Mathematically;

$$\sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

- Squashes real numbers to range between  $[-1, 1]$ .



- Scaled version of sigmoid neuron.

$$\tanh(x) = \frac{2}{1+e^{-2x}} - 1$$

### Pros and Cons.


- Outperforms the sigmoid activation functions.
- Zero centered output.



## 1.2 Some Common Non Linear Activation Functions.

### ReLU


- It computes the function  $f(x)=\max(0,x)$ .



- Pros:
  - It was found to greatly accelerate the convergence of (stochastic) gradient descent compared to the sigmoid/tanh functions.
  - Compared to tanh/sigmoid neurons that involve expensive operations (exponentials, etc.), the ReLU can be implemented by simply thresholding a matrix of activations at zero.
- Cons:
  - ReLU units can be fragile during training and can “die”.
    - What if, the current weights put the ReLU on the left flat side while it optimally should be on the right side for this particular input ?
    - The gradient is 0 and so the weight will not be updated, not even a tiny bit, so where is "learning" in this case?

### Leaky ReLU

- Leaky ReLUs are one attempt to fix the “dying ReLU” problem.
- That is, the function computes:
$$f(x)=1(x<0)(\alpha x)+1(x\geq 0)(x)$$
- where  $\alpha$  is a small constant.



## 1.3 What activation should I use?

- Never Use Sigmoid.
- Try tanh, but expect it to work worse than ReLU.
- While using ReLU/Leaky ReLU, be sensible while picking the learning rates.
- If possible, monitor the dead neuron in a network.


## **2. Error/Loss Function.**

## 2.1 Cross Entropy Loss.

- CEL measures the distance between two probability distributions.

$$H(P^* | P) = - \sum_i P^*(i) \log P(i)$$

- In information theory, information are quantified using the level of surprise any event can cause.
  - $h(x) = -\log(p(x))$
  - Low Probability Event: High Information (surprising).
  - High Probability Event: Low Information (unsurprising).



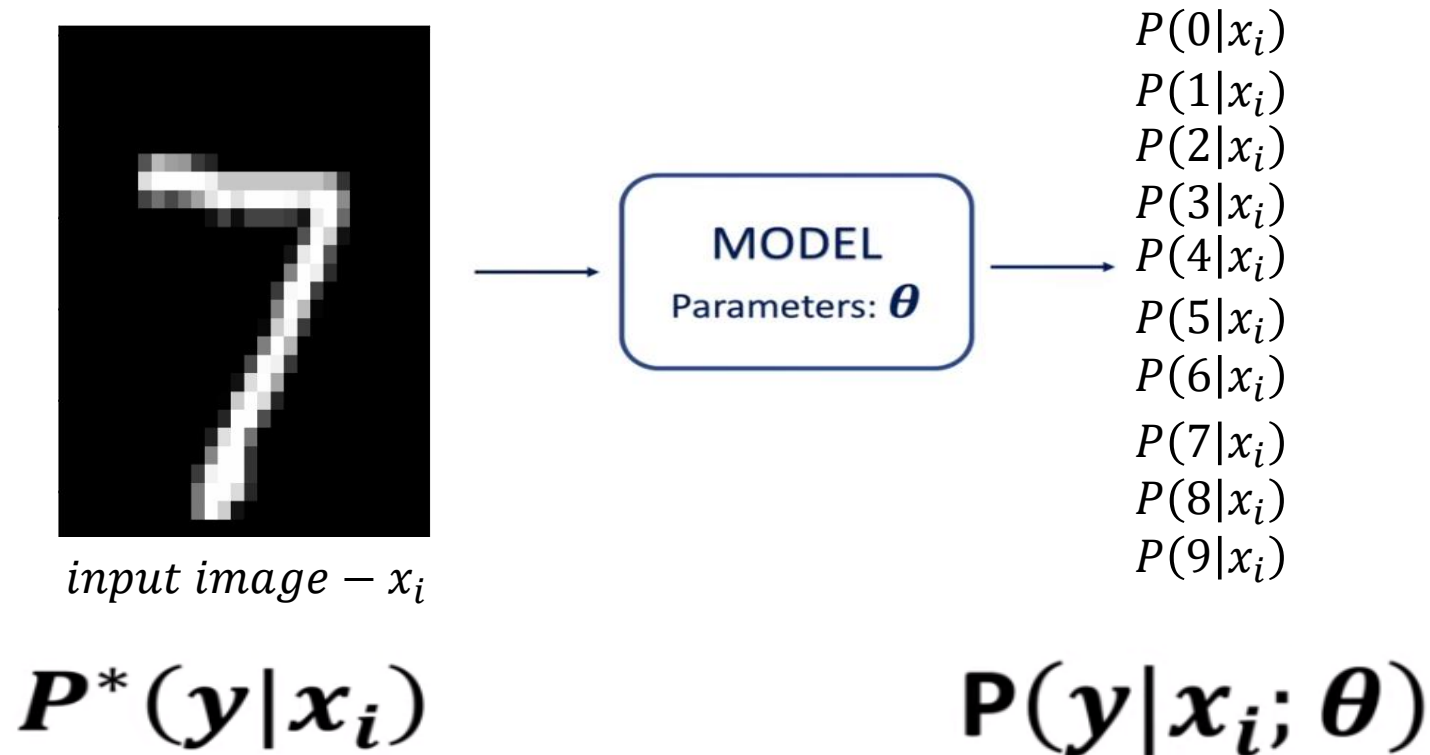
Where are We?

1.Activation Functions.

2.Error/Loss Functions.

2.1 Cross Entropy Loss.

## 2.2 Problem Setup-Classification.




$$\begin{aligned} D_{KL}(\overbrace{P^*(y|x_i)}^{\text{TRUE CLASS DISTRIBUTION}} || \overbrace{P(y|x_i; \theta)}^{\text{PREDICTED CLASS DISTRIBUTION}}) \\ D_{KL}(P^*||P) &= \sum_y P^*(y|x_i) \log \frac{P^*(y|x_i)}{P(y|x_i; \theta)} \\ &= \sum_y P^*(y|x_i) [\log P^*(y|x_i) - \log P(y|x_i; \theta)] \\ &= \sum_y \overbrace{P^*(y|x_i) \log P^*(y|x_i)}^{\text{DOESN'T DEPEND ON } \theta} - \sum_y P^*(y|x_i) \log P(y|x_i; \theta) \\ \operatorname{argmin}_{\theta} D_{KL}(P^*||P) &\equiv \operatorname{argmin}_{\theta} - \sum_y P^*(y|x_i) \log P(y|x_i; \theta) \end{aligned}$$

# 3. Calculation of Gradient.


How do we Learn Weights?

# 3.1 Forward and Backward Propagations.

- The weights in Multi layer networks are learned with the combinations of forward and backward propagations.
- a network forward propagates activation to produce an output and it backward propagates error to determine weight changes



## 3.1 Feed Forward Calculations.



$$Z^{[1]} = W^{[1]}X + b^{[1]}$$

$$A^{[1]} = g^{[1]}(Z^{[1]})$$

$$Z^{[2]} = W^{[2]}A^{[1]} + b^{[2]}$$

$$A^{[2]} = g^{[2]}(Z^{[2]})$$


$$\vdots$$

$$A^{[L]} = g^{[L]}(Z^{[L]}) = \hat{Y}$$



## 3.2 Back-Propagation.

- Backpropagation is a technique used by deep layer networks to find the error of the network.
- The error is calculated by comparing an expected output with a predicted output, this algorithm then propagated these errors backward to update weights and biases.




# Backpropagation Intuition and Computational Graph.

# Computational Graph -1

$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$



# Computational Graph -2


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Computational Graph -3


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial f}$$

# Computational Graph -4


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Computational Graph -5


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Computational Graph -6


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$



# Computational Graph -7


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Computational Graph -8


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Computational Graph -9


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Computational Graph -10


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial y}$$

# Computational Graph -11


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

$$\frac{\partial f}{\partial y}$$

# Computational Graph -12


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Computational Graph -13


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$




Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

# Backpropagation Requirements

- Backprop Requires following three things:
  1. Dataset: in the pairs of input-output i.e.  $(x_i, y_i)$ .
  2. A feed-forward neural network;
  3. An Error Function, E which defines the error between the desired output and calculated output.





# Backpropagation Algorithm.

---

**Algorithm 2** Backpropagation algorithm for feedforward networks.

---

**Require:** a set of training examples  $D$ , learning rate  $\alpha$ .

1. Create a feed-forward network with  $n_{in}$  inputs,  $n_h$  hidden units, and  $n_o$  output units.
2. Initialize all network weights to *small* random numbers.
3. **Repeat** until the termination condition is met:

**For** each training example  $(\mathbf{x}, \mathbf{t}) \in D$  **do**:

*Propagate* the input  $\mathbf{x}$  forward through the network, i.e.:

1. Input  $\mathbf{x}$  to the network and compute the output  $a_k$  of units in the output layer.

*Backpropagate* the errors through the network:

1. **For** each network output unit  $k$ , calculate its error term  $\delta_k$ :

$$\delta_k \leftarrow -a_k(1 - a_k)(t_k - a_k) \quad (3)$$

2. **For** each hidden unit  $h$ , calculate its error term  $\delta_h$ :

$$\delta_h \leftarrow a_h(1 - a_h) \sum_k \delta_k w_{kh} \quad (4)$$

3. Update each network weight  $w_{ji}$

$$w_{ji} \leftarrow w_{ji} + \Delta w_{ji}$$

where

$$\Delta w_{ji} = -\alpha \delta_j x_{ji}$$

---

Notations:

The subscript  $k$  denotes the output layer.

The subscript  $j$  denotes the hidden layer.

The subscript  $i$  denotes the input layer


# In Tutorial.

- Pen paper calculation on Forward and Backward Propagation.

# 4. Optimization in Deep Learning.


# 4.1 Optimization.

- Optimization and machine learning have related, but somewhat different goals
  - Goal in optimization: minimize an objective function
    - For a set of training examples, reduce the **training error**
  - Goal in ML: find a suitable model, to predict on data examples
    - For a set of testing examples, reduce the **generalization error**
- For a given empirical function  $g$  (dashed purple curve), optimization algorithms attempt to find the point of minimum **empirical risk**
- The expected function  $f$  (blue curve) is obtained given a limited amount of training data examples
- ML algorithms attempt to find the point of minimum **expected risk**, based on minimizing the error on a set of testing examples
  - Which may be at a different location than the minimum of the training examples
  - And which may not be minimal in a formal sense




## 4.2 Stationary points.

- **Stationary points** ( or **critical points**) of a differentiable function  $f(x)$  of one variable are the points where the derivative of the function is zero, i.e.,  $f'(x) = 0$
- The stationary points can be:
  - **Minimum**, a point where the derivative changes from negative to positive
  - **Maximum**, a point where the derivative changes from positive to negative
  - **Saddle point**, derivative is either positive or negative on both sides of the point
- The minimum and maximum points are collectively known as **extremum points**
- The nature of stationary points can be determined based on the second derivative of  $f(x)$  at the point
  - If  $f''(x) > 0$ , the point is a minimum
  - If  $f''(x) < 0$ , the point is a maximum
  - If  $f''(x) = 0$ , inconclusive, the point can be a saddle point, but it may not
- The same concept also applies to gradients of multivariate functions





## 4.3 Local Minima.

- Among the challenges in optimization of model's parameters in ML involve local minima, saddle points, vanishing gradients
- For an objective function  $f(x)$ , if the value at a point  $x$  is the minimum of the objective function **over the entire domain** of  $x$ , then it is the *global minimum*
- If the value of  $f(x)$  at  $x$  is smaller than the values of the objective function at any other points in **the vicinity** of  $x$ , then it is the *local minimum*
  - The objective functions in ML usually have many local minima
    - When the solution of the optimization algorithm is near the local minimum, the gradient of the loss function approaches or becomes zero (vanishing gradients)
    - Therefore, the obtained solution in the final iteration can be a local minimum, rather than the global minimum



## 4.4 Saddle Points.

- The gradient of a function  $f(x)$  at a **saddle point** is 0, but the point is not a minimum or maximum point
  - The optimization algorithms may stall at saddle points, without reaching a minima
- Note also that the point of a function at which the sign of the curvature changes is called an **inflection point**
  - An inflection point ( $f''(x) = 0$ ) can also be a saddle point, but it does not have to be
- For the 2D function (right figure), the saddle point is at (0,0)
  - The point looks like a saddle, and gives the minimum with respect to  $x$ , and the maximum with respect to  $y$






# 5. Optimization via Gradient Descent.



# 5.1 Pseudo-code Gradient Descent.

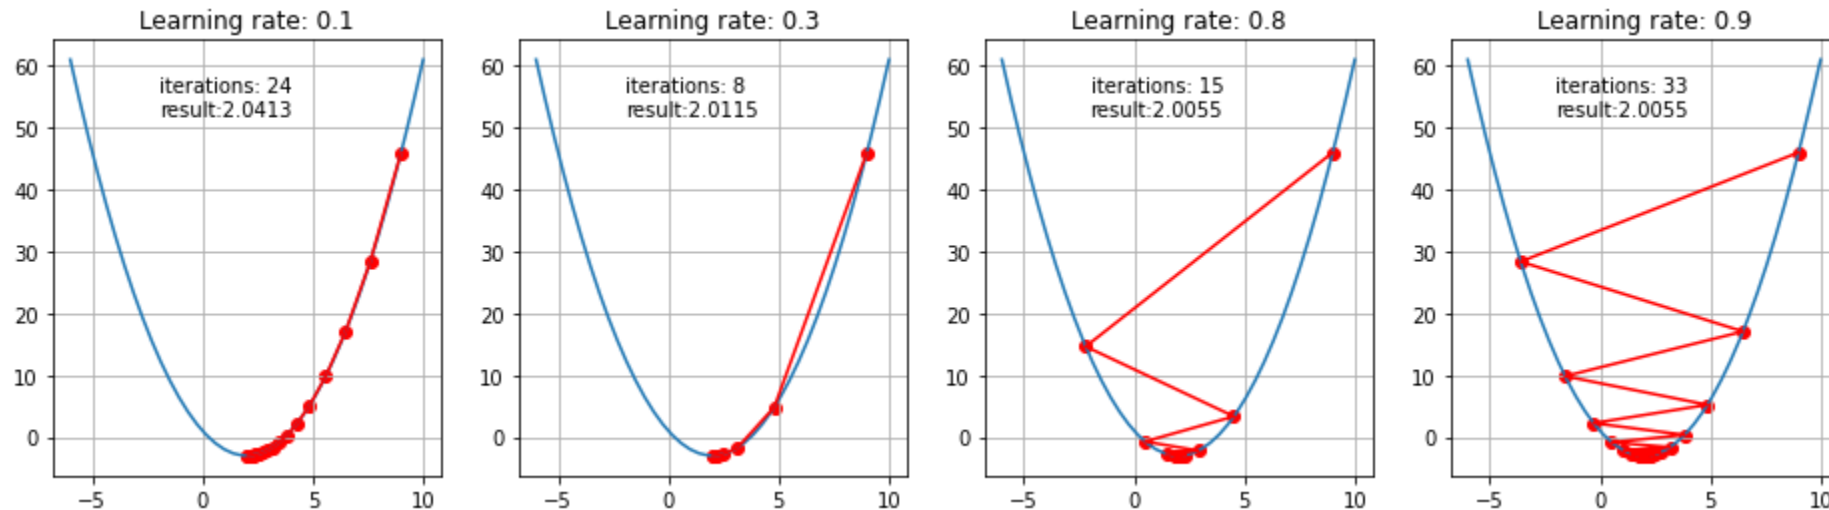
- Choose a starting point.(Initialization)
- Calculate the gradient at this point.
- Make a scaled step in the opposite direction to the gradient.
- Repeat until stopping criteria is met, which can be:
  - Maximum number of iterations reached.
  - Step size is smaller than tolerance.
- Input Parameters:
  - Starting point: 0 or random initialization.
  - Gradient function: Cost function to be minimized.
  - Learning Rate: scaling factor for step sizes.
  - Max. Iterations: number of iterations the algorithm must run.

# 5.2 Gradient Descent: An Illustration




## 5.3 Learning Rate and Convergence.

- High Learning Rate does not Guarantees speedy Convergence.




## 5.4 Starting point and Convergence.

- Convergence speed also depends on our starting point.



## 5.5 Gradient descent in practice: Learning rate

- Automatic convergence test
- $\alpha$  too small: slow convergence
- $\alpha$  too large: may not converge
- To choose  $\alpha$ , try 0.001, ... 0.01, ..., 0.1, ... , 1



# **6. Variants of Gradient Decent.**

# 6.1 Gradient descent variants.

- Based on how much data we use to compute the gradient of the objective function, there are three main variants of gradient descent:
  - Batch Gradient Descent.
  - Stochastic Gradient Descent.
  - Mini-Batch Gradient Descent.
- More data may mean more accuracy of the parameter update, but it also means more time needed to reach convergence, thus the **trade off** is necessity.


## 6.2 Batch Gradient Descents.

- Also known as vanilla gradient descent, computes the gradient of the cost function w.r.t to the parameters  $\theta$  for the entire training set:
  - **Update rule:**
    - $\theta = \theta - \alpha \nabla_{\theta} J(\theta)$ .
- It can be very slow and is intractable for datasets that do not fit the memory.



## 6.3 Stochastic Gradient Descent.

- In contrast SGD performs a parameter update for each training example i.e.  $x^i$  and label  $y^i$ .
  - **Update Rule:**
  - $\theta = \theta - \alpha \nabla_{\theta} J(\theta; x^i; y^i)$ .
- It is usually much faster compared to BGD, and also can be used to learn online.
- SGD performs frequent updates with a high variance that cause the objective function to fluctuate heavily as in image:
- It can enable it to jump to new and potential better local minima.
- It can also ultimately complicates the convergence to the exact minimum.



## 6.4 Mini Batch Gradient Descent.

- It is the mixture of BGD and SGD i.e. it updates the parameter for every mini batch of  $n$  training examples:
  - **Update Rule:**
    - $\theta = \theta - \alpha \nabla_{\theta} J(\theta; \mathbf{x}^{(i:i+n)}; \mathbf{y}^{(i:i+n)})$ .
- Merits:
  - Reduces the variance of the parameter updates, which can lead to more stable convergence;
  - Efficient computation for large models.
- Common mini-batch size range between 50 and 256.

# 6.5 Challenges


- Same learning rate applies to all parameter updates.
- Getting trapped into suboptimal local minima or saddle points.
- Choosing a Proper Learning Rate.
  - Annealing:
    - Learning rate schedules or reducing the learning rate according to a pre-defined schedule or it becomes smaller than a pre-set threshold.
    - Schedules and threshold values have to be pre-defined, thus may not be able to adapt to a dataset's characteristics.

# 6.6 Basic Mini-Batch Gradient Descent.

- Guess an **initial learning rate**.
  - If the **error** keeps getting worse or **oscillates** wildly, **reduce** the **learning rate**.
  - If the **error** is falling **fairly** consistently but slowly, increase the **learning rate**.
- Write a simple program to automate this way of **adjusting** the **learning rate**.
- Towards the end of mini-batch learning it nearly always helps to **turn down** the **learning rate**.
  - This removes **fluctuations** in the **final weights** caused by the **variations between mini-batches**.
- Turn down the **learning rate** when the **error stops decreasing**.
  - Use the error on a separate validation set

## 6.7 Turning Down the learning rate.

- Turning down the learning rate reduces the random fluctuations in the error due to the different gradients on different mini-batches.
  - So we get a quick win.
  - But then we get slower learning.
- Don't turn down the learning rate too soon!



# Today's Lesson...

- Activations Function.
- Error Function in Classification with Neural Network.
- Learn the weights:
- Forward Pass.
- Backward Pass.
- Optimization-Gradient Descent and its Variants.

## Next Up!!!!

- You are given a training set with 1M labeled points. When you train a shallow neural net with one fully-connected feed-forward hidden layer on this data you obtain 86% accuracy on test data. When you train a deeper neural net as in which consist of a convolutional layer, pooling layer, and three fully-connected feed-forward layers on the same data you obtain 91% accuracy on the same test set.
- What is the source of this improvement?

# At the end.....

- Questions?



- **Reminder!!!!**
  - **Assignment-I is out.**
  - **Date to submit-26 March.**